

OBJECTIVE

1. APPLYING RANDOM FOREST WITH TFIDF VECTORIZATION

- FINDING THE BEST HYPERPARAMETER USING GRIDSEARCHCV WITH TRAIN DATA AND CROSS-VALIDATION DATA BY PLOTTING THE RESULTS OF VARIOUS TRAIN DATA AND CROSS VALIDATION DATA
- USING THE APPROPRIATE VALUE OF HYPERPARAMETER, TESTING ACCURACY ON TEST DATA USING F1-SCORE
- PLOTTING THE CONFUSION MATRIX TO GET THE PRECISION, RECALL VALUE WITH HELP OF HEATMAP
- PRINTING THE TOP 30 MOST IMPORTANT FEATURES

```
In [0]: from sklearn.model_selection import train_test_split          #importin
        g the necessary libraries
        from sklearn.model_selection import RandomizedSearchCV
        from sklearn.datasets import *
        from sklearn import naive_bayes
        from sklearn.feature_extraction.text import CountVectorizer
        from sklearn.feature_extraction.text import TfidfVectorizer
        import numpy as np
        import pandas as pd
        from sklearn import *
        import warnings
        warnings.filterwarnings("ignore")
        from sklearn.ensemble import RandomForestClassifier
```

```
In [4]: from google.colab import drive
        drive.mount('/content/gdrive')#getting the content from the google driv
        e
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

```

In [0]: final_processed_data=pd.read_csv("gdrive/My Drive/final_new_data.csv")#
        loading the preprocessed data with 100k points into dataframe

In [6]: # getting the counts of 0 and 1 in "SCORE" column to know whether it is
        unbalanced data or not
        count_of_1=0
        count_of_0=0
        for i in final_processed_data['Score']:
            if i==1:
                count_of_1+=1
            else:
                count_of_0+=1
        print(count_of_1)
        print(count_of_0)
        #it is an imbalanced dataset

88521
11479

In [0]: #spliiting the data into train and test data
        x_train,x_test,y_train,y_test=model_selection.train_test_split(final_pr
        ocessed_data['CleanedText'].values,final_processed_data['Score'].values
        ,test_size=0.3,shuffle=False)

In [8]: vectorizer=TfidfVectorizer(min_df=2)#building the vertorizer with word
        counts equal and more then 2
        train_tfidf=vectorizer.fit_transform(x_train)#fitting the model on trai
        ning data
        print(train_tfidf.shape)

(70000, 16382)

In [9]: test_tfidf=vectorizer.transform(x_test)#fitting the bow model on test d
        ata
        print("shape of x_test after tfidf vectorization ",test_tfidf.shape)

shape of x_test after tfidf vectorization (30000, 16382)

```

In [0]:

```
In [10]: #building the model  
#using time series split method for cross-validation score  
from sklearn.model_selection import TimeSeriesSplit  
tscv = TimeSeriesSplit(n_splits=5)  
rf=RandomForestClassifier(criterion='gini',class_weight={1:.5,0:.5})  
tuned_parameters=[{'max_depth':[61,64,68,73,77,80], 'n_estimators':[21,30,35,40,45,50]}]  
#applying the model of decision tree and using gridsearchcv to find the best hyper parameter  
%%time  
from sklearn.model_selection import GridSearchCV  
model = GridSearchCV(rf, tuned_parameters, scoring = 'f1', cv=tscv,n_jobs=-1)#building the gridsearchcv model
```

CPU times: user 4 µs, sys: 1 µs, total: 5 µs
Wall time: 7.63 µs

```
In [11]: %%time  
model.fit(train_tfidf, y_train)#fitting the training data
```

CPU times: user 11.2 s, sys: 284 ms, total: 11.5 s
Wall time: 26min 40s

```
Out[11]: GridSearchCV(cv=TimeSeriesSplit(max_train_size=None, n_splits=5),  
                      error_score='raise-deprecating',  
                      estimator=RandomForestClassifier(bootstrap=True, class_weight=  
{1: 0.5, 0: 0.5},  
                      criterion='gini', max_depth=None, max_features='auto',  
                      max_leaf_nodes=None, min_impurity_decrease=0.0,  
                      min_impurity_split=None, min_samples_leaf=1,  
                      min_samples_split=2, min_weight_fraction_leaf=0.0,  
                      n_estimators='warn', n_jobs=None, oob_score=False,  
                      random_state=None, verbose=0, warm_start=False),  
                      fit_params=None, iid='warn', n_jobs=-1,  
                      param_grid=[{'max_depth': [61, 64, 68, 73, 77, 80], 'n_estimator  
s': [21, 30, 35, 40, 45, 50]}],
```

```
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring='f1', verbose=0)
```

In [20]: `print(model.best_estimator_)`*#printing the best_estimator*

```
RandomForestClassifier(bootstrap=True, class_weight={1: 0.5, 0: 0.5},
                        criterion='gini', max_depth=73, max_features='auto',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=21, n_jobs=None, oob_score=False,
                        random_state=None, verbose=0, warm_start=False)
```

In [21]: `print(model.score(test_tfidf,y_test))`*#checking the score on test_Data*

```
0.9376672494380418
```

In [22]: `results=pd.DataFrame(model.cv_results_)`*# getting varoius cv_scores and train_scores various values of hyperparameter given as parameter and storing it in a dataframe*
`results`*#printing the dataframe*

Out[22]:

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_max_d
0	8.207892	0.125787	0.942826	0.968657	61
1	11.751353	0.175003	0.942623	0.968801	61

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_max_d
2	13.598304	0.196765	0.942794	0.968604	61
3	15.517566	0.228227	0.942471	0.968346	61
4	17.657933	0.251955	0.942571	0.968759	61
5	19.466890	0.272821	0.942472	0.969249	61
6	8.745539	0.132947	0.942786	0.970711	64
7	12.032980	0.176906	0.943092	0.970263	64
8	14.123110	0.203084	0.942750	0.970345	64


	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_max_d
9	16.386997	0.228492	0.942637	0.970659	64
10	18.190747	0.255849	0.942665	0.970313	64
11	20.558136	0.284757	0.942728	0.970934	64
12	9.144402	0.134296	0.943007	0.972365	68
13	12.795353	0.185571	0.942633	0.971862	68
14	15.200046	0.212584	0.942648	0.972416	68
15	17.023168	0.234259	0.942667	0.972219	68

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_max_d
16	19.363737	0.266417	0.942464	0.972856	68
17	21.581563	0.289683	0.942622	0.971730	68
18	9.801576	0.143186	0.943570	0.974136	73
19	14.160118	0.194498	0.943043	0.973793	73
20	16.824344	0.218368	0.943022	0.974422	73
21	18.932823	0.247959	0.943131	0.973988	73
22	21.523917	0.281439	0.942929	0.974456	73

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_max_d
23	24.093993	0.315355	0.942896	0.974129	73
24	10.646736	0.146012	0.943391	0.975058	77
25	15.027038	0.198350	0.943395	0.975380	77
26	17.612787	0.228716	0.943262	0.976097	77
27	19.867350	0.254089	0.943224	0.975760	77
28	22.579286	0.295588	0.943239	0.975469	77
29	24.659954	0.314257	0.943005	0.976081	77

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_max_d
30	10.896179	0.148836	0.943315	0.976295	80
31	15.231942	0.198558	0.943331	0.976448	80
32	17.803722	0.228328	0.943447	0.976752	80
33	20.242478	0.261301	0.943186	0.977049	80
34	22.851599	0.298497	0.943296	0.976556	80
35	23.545404	0.307550	0.943202	0.977062	80

36 rows × 22 columns



```
In [0]: results['mean_test_score']=results['mean_test_score']*100
results=results.round(decimals=2)
```

```
results['cv_error_score']=100-results['mean_test_score']
```

PLOTTING THE HEATMAP WITH HYPERPARAMETERS FOR CV_ERROR SCORE

```
In [0]: test_score_heatmap=results.pivot(          'param_max_depth'          , 'param  
_n_estimators', 'cv_error_score' )
```

```
In [25]: test_score_heatmap
```

Out[25]:

param_n_estimators	21	30	35	40	45	50
param_max_depth						
61	5.72	5.74	5.72	5.75	5.74	5.75
64	5.72	5.69	5.72	5.74	5.73	5.73
68	5.70	5.74	5.74	5.73	5.75	5.74
73	5.64	5.70	5.70	5.69	5.71	5.71
77	5.66	5.66	5.67	5.68	5.68	5.70
80	5.67	5.67	5.66	5.68	5.67	5.68

```
In [26]: import seaborn as sns  
sns.heatmap(test_score_heatmap,annot=True,annot_kws={"size": 15}, fmt=  
'g',linewidths=.3)
```

Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb8be172a58>



In [27]: `print(model.best_estimator_)`*#printing the best_estimator*


```
RandomForestClassifier(bootstrap=True, class_weight={1: 0.5, 0: 0.5},
                        criterion='gini', max_depth=73, max_features='auto',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=21, n_jobs=None, oob_score=False,
                        random_state=None, verbose=0, warm_start=False)
```

FROM THE ABOVE HEATMAPS RESULTS FOR CV DATA, WE FOUND THAT BEST HYPERPARAMETERS AS MAX_DEPTH=73 AND N_ESTIMATORS=21

PLOTTING THE ROC CURVE FOR GETTING AUC SCORE

```
In [29]: rf=RandomForestClassifier(criterion='gini',class_weight={1:.5,0:.5},max
_depth=73 ,n_estimators=21)
rf.fit(train_tfidf,y_train)#fitting the model
probs = rf.predict_proba(test_tfidf)
preds = probs[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

#
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'best')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
In [30]: print('accuracy from the ROC curve is found as ',roc_auc*100)
accuracy from the ROC curve is found as 85.30101633953345
```

```
In [31]: z=rf.feature_importances_
a=z.argsort()
print('shape of wieght vector is:',a.shape)
top_features=np.take(vectorizer.get_feature_names(),a[16360:])#taking 1
ast features as they are of very high importance

shape of wieght vector is: (16382,)
```

```
In [32]: print(top_features)#printing the top_features
top=list(top_features)

['return' 'refund' 'receiv' 'poor' 'gross' 'money' 'best' 'thought' 'ta
st']
```

```
'horribl' 'mayb' 'stale' 'love' 'threw' 'wast' 'bad' 'worst' 'terribl'  
'aw' 'would' 'great' 'disappoint']
```

REPRESENTING TOP IMPORTANT FEATURES USING WORDCLOUD LIBRARY

```
In [33]: from wordcloud import WordCloud #here we are printing the top features  
         using wordcloud library  
         import matplotlib.pyplot as plt  
         wordcloud = WordCloud(width = 1500, height = 1000,  
                                background_color = 'white',  
                                min_font_size = 10).generate(str(top))  
  
         # plot the WordCloud image  
         plt.figure(figsize = (8, 8), facecolor = None)  
         plt.imshow(wordcloud)  
         plt.axis("off")  
         plt.tight_layout(pad = 0)  
  
         plt.show()
```




TESTING OUR MODEL ON TEST DATA AND CHECKING ITS PRECISION ,RECALL ,F1_FCORE

```
In [34]: #Testing Accuracy on Test data
import seaborn as sns #importing seaborn as sns
from sklearn.metrics import *#importing varoius metrics from sklearn
#building the model
y_pred = rf.predict(test_tfidf)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100))#printing accuracy
print("Precision on test set: %0.3f"%(precision_score(y_test, y_pred)))
```

```
#printing precision score
print("Recall on test set: %0.3f"%(recall_score(y_test, y_pred))) #printing recall
print("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred)))
print("Confusion Matrix of test set:\n [ [TN  FP]\n [FN TP] ]\n")
df_cm = pd.DataFrame(confusion_matrix(y_test, y_pred), range(2),range(2)) #generating the heatmap for confusion matrix
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g',cmap="YlGnBu")
```

```
Accuracy on test set: 88.443%
Precision on test set: 0.884
Recall on test set: 0.999
F1-Score on test set: 0.938
Confusion Matrix of test set:
 [ [TN  FP]
   [FN TP] ]
```

Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb8b91d6e80>



BOW VECTORIZATION FOR RANDOM FOREST IS COMPLETED

OBJECTIVE

1. APPLYING GBDT WITH TFIDF VECTORIZATION

- FINDING THE BEST HYPERPARAMETER USING GRIDSEARCHCV WITH TRAIN DATA AND CROSS-VALIDATION DATA BY PLOTTING THE RESULTS OF VARIOUS TRAIN DATA AND CROSS VALIDATION DATA
- USING THE APPROPRIATE VALUE OF HYPERPARAMETER, TESTING ACCURACY ON TEST DATA USING F1-SCORE
- PLOTTING THE CONFUSION MATRIX TO GET THE PRECISION, RECALL VALUE WITH HELP OF HEATMAP

- PRINTING THE TOP 30 MOST IMPORTANT FEATURES

```
In [35]: from xgboost import XGBClassifier
#building the model
#using time series split method for cross-validation score
from sklearn.model_selection import TimeSeriesSplit
tscv = TimeSeriesSplit(n_splits=5)
xg=XGBClassifier(n_jobs=-1)
tuned_parameters=[{'max_depth': [11,15,20,24,27,30], 'n_estimators': [21,30,35,40,45,50]}]
#applying the model of decision tree and using gridsearchcv to find the best hyper parameter
%%time
from sklearn.model_selection import GridSearchCV
model = GridSearchCV(xg, tuned_parameters, scoring = 'f1', cv=tscv,n_jobs=-1)#building the gridsearchcv model
```

CPU times: user 3 µs, sys: 1 µs, total: 4 µs
Wall time: 6.91 µs

```
In [36]: %%time
model.fit(train_tfidf, y_train)#fitting the training data
```

CPU times: user 2min 51s, sys: 362 ms, total: 2min 51s
Wall time: 1h 12min 21s

```
Out[36]: GridSearchCV(cv=TimeSeriesSplit(max_train_size=None, n_splits=5),
                    error_score='raise-deprecating',
                    estimator=XGBClassifier(base_score=0.5, booster='gbtree', colsam
ple_bylevel=1,
                    colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=
0,
                    max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
n_jobs=-1, nthread=None, objective='binary:logistic',
                    random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
                    seed=None, silent=True, subsample=1),
                    fit_params=None, iid='warn', n_jobs=-1,
                    param_grid=[{'max_depth': [11, 15, 20, 24, 27, 30], 'n_estimator
s': [21, 30, 35, 40, 45, 50]}],
```

```
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring='f1', verbose=0)
```

```
In [43]: print(model.best_estimator_)#printing the best_estimator
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=
0,
              max_depth=30, min_child_weight=1, missing=None, n_estimators=50,
              n_jobs=-1, nthread=None, objective='binary:logistic',
              random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
              seed=None, silent=True, subsample=1)
```


```
In [60]: results=pd.DataFrame(model.cv_results_)# getting various cv_scores and
         train_scores various values of hyperparameter given as parameter and s
         toring it in a dataframe
         results.head()#printing the dataframe
```

Out[60]:

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_max_de
0	14.193161	0.189327	0.946074	0.959469	11
1	20.153732	0.204038	0.947127	0.962851	11
2	23.310932	0.210168	0.947627	0.964501	11

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_max_de
3	26.559028	0.217920	0.947915	0.966236	11
4	29.652577	0.224765	0.948380	0.967831	11

5 rows × 22 columns



```
In [0]: results['mean_test_score']=results['mean_test_score']*100
results=results.round(decimals=2)
results['cv_error_score']=100-results['mean_test_score']
```

PLOTTING THE HEATMAP WITH HYPERPARAMETERS FOR CV_ERROR SCORE

```
In [48]: test_score_heatmap=results.pivot( 'param_max_depth' , 'param_n_estimator
s' , 'cv_error_score' )
test_score_heatmap
```

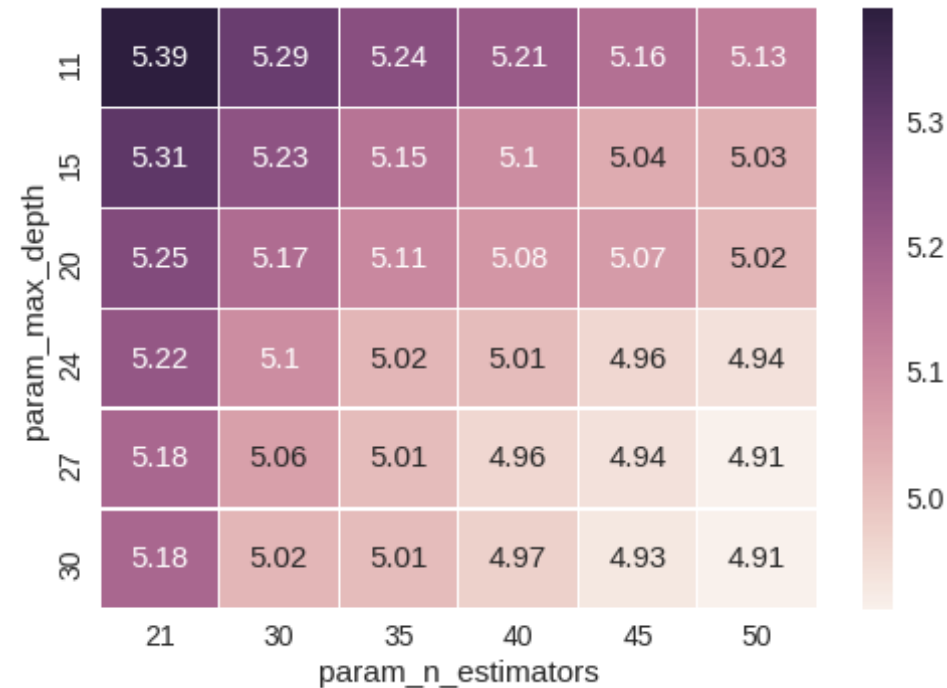
Out[48]:

param_n_estimators	21	30	35	40	45	50
param_max_depth						

param_n_estimators	21	30	35	40	45	50
param_max_depth						
11	5.39	5.29	5.24	5.21	5.16	5.13
15	5.31	5.23	5.15	5.10	5.04	5.03
20	5.25	5.17	5.11	5.08	5.07	5.02
24	5.22	5.10	5.02	5.01	4.96	4.94
27	5.18	5.06	5.01	4.96	4.94	4.91
30	5.18	5.02	5.01	4.97	4.93	4.91

```
In [49]: import seaborn as sns
sns.heatmap(test_score_heatmap,annot=True,annot_kws={"size": 15}, fmt=
'g',linewidths=.3)
```

```
Out[49]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb8b81160b8>
```



```
In [50]: print(model.best_estimator_)#printing the best_estimator
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bytree=1, gamma=0, learning_rate=0.1, max_delta_step=
0,
              max_depth=30, min_child_weight=1, missing=None, n_estimators=50,
              n_jobs=-1, nthread=None, objective='binary:logistic',
              random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
              seed=None, silent=True, subsample=1)
```


FROM THE ABOVE HEATMAPS RESULTS FOR CV DATA,WE FOUND THAT BEST HYPERPARAMETERS AS MAX_DEPTH=30 AND N_ESTIMATORS=50

PLOTTING THE ROC CURVE FOR GETTING AUC SCORE

```
In [0]: xg=XGBClassifier(n_jobs=-1,max_depth=30 ,n_estimators=50)
xg.fit(train_tfidf,y_train)#fitting the model
```

```
In [53]: probs = xg.predict_proba(test_tfidf)
preds = probs[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'best')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
In [54]: print('accuracy from the ROC curve is found as ',roc_auc*100)
```

accuracy from the ROC curve is found as 87.37248689754394

```
In [55]: z=rf.feature_importances_
a=z.argsort()
print('shape of wieght vector is:',a.shape)
top_features=np.take(vectorizer.get_feature_names(),a[16360:])#taking 1
last features as they are of very high importance
```

shape of wieght vector is: (16382,)


```
In [56]: print(top_features)#printing the top_features
top=list(top_features)
```

['return' 'refund' 'receiv' 'poor' 'gross' 'money' 'best' 'thought' 'ta


```
st'  
'horribl' 'mayb' 'stale' 'love' 'threw' 'wast' 'bad' 'worst' 'terribl'  
'aw' 'would' 'great' 'disappoint']
```

REPRESENTING TOP IMPORTANT FEATURES USING WORDCLOUD LIBRARY

```
In [57]: from wordcloud import WordCloud #here we are printing the top features  
         using wordcloud library  
         import matplotlib.pyplot as plt  
         wordcloud = WordCloud(width = 1500, height = 1000,  
                                background_color = 'white',  
                                min_font_size = 10).generate(str(top))  
  
         # plot the WordCloud image  
         plt.figure(figsize = (8, 8), facecolor = None)  
         plt.imshow(wordcloud)  
         plt.axis("off")  
         plt.tight_layout(pad = 0)  
  
         plt.show()
```



```
In [59]: y_pred = xg.predict(test_tfidf)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100))#printing accuracy
print("Precision on test set: %0.3f"%(precision_score(y_test, y_pred)))
#printing precision score
print("Recall on test set: %0.3f"%(recall_score(y_test, y_pred))) #printing recall
print("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred)))
print("Confusion Matrix of test set:\n [ [TN FP]\n [FN TP] ]\n")
df_cm = pd.DataFrame(confusion_matrix(y_test, y_pred), range(2),range(2)) #generating the heatmap for confusion matrix
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Accuracy on test set: 90.110%

```
Precision on test set: 0.905
Recall on test set: 0.991
F1-Score on test set: 0.946
Confusion Matrix of test set:
[ [TN FP]
  [FN TP] ]
```

Out[59]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb8b7f11128>

