

OBJECTIVE :

1. APPLYING LOGISTIC REGRESSION WITH BOW VECTORIZATION

- PERFORMING PERTUBATION TEST TO CHECK WHETHER OUR DATA FEATURES ARE COLLINER OR NOT AND PLOTTING THE RESULT
- FINDING THE BEST HYPERPARAMETER USING GRIDSEARCHCV WITH TRAIN DATA AND CROSS-VALIDATION DATA BY PLOTTING THE RESLUTS OF VAROIUS TRAIN DATA AND CROSS VALIDATION DATA
- USING THE APROPRIATE VALUE OF HYPERPARAMETER ,TESTING ACCURACY ON TEST DATA USING F1-SCORE
- PLOTTING THE CONFUSION MATRIX TO GET THE PRECISOIN ,RECALL VALUE WITH HELP OF HEATMAP
- PRINTING THE TOP 20 FEATURES FOR BOTH POSITIVE AND NEGATIVE WORDS #

```
In [2]: from sklearn.model_selection import train_test_split          #importin
        g the necessary libraries
        from sklearn.model_selection import RandomizedSearchCV
        from sklearn.datasets import *
        from sklearn import naive_bayes
        from sklearn.feature_extraction.text import CountVectorizer
        from sklearn.feature_extraction.text import TfidfVectorizer
        import numpy as np
        import pandas as pd
        from sklearn import *
        import warnings
        warnings.filterwarnings("ignore")
```

```
In [3]: final_processed_data=pd.read_csv("C:/Users/Mayank/Desktop/machine learn
        ing/applieaiaicourse data/lecture 18 knn/final_new_data.csv")#loading t
        he preprocessed data with 100k points into dataframe
```

```
In [4]: # getting the counts of 0 and 1 in "SCORE" column to know whether it is
        # unbalanced data or not
        count_of_1=0
        count_of_0=0
        for i in final_processed_data['Score']:
            if i==1:
                count_of_1+=1
            else:
                count_of_0+=1
        print(count_of_1)
        print(count_of_0)
        #it is an imbalanced dataset

88521
11479
```

```
In [5]: #splitting the data into train and test data
        x_train,x_test,y_train,y_test=model_selection.train_test_split(final_pr
        ocessed_data['CleanedText'].values,final_processed_data['Score'].values
        ,test_size=0.2,shuffle=False)
```

```
In [6]: vectorizer=CountVectorizer(min_df=10)#building the vectorizer with word
        #counts equal and more then 2
        train_bow=vectorizer.fit_transform(x_train)#fitting the model on traini
        ng data
        print(train_bow.shape)

(80000, 7677)
```

```
In [7]: from sklearn.preprocessing import StandardScaler #standarizing the trai
        ning data
        x_train_data=StandardScaler( with_mean=False).fit_transform(train_bow)
        print(x_train_data.shape)

(80000, 7677)
```

```
In [8]: test_bow=vectorizer.transform(x_test)#fitting the bow model on test dat
        a
        print("shape of x_test after bow vectorization ",test_bow.shape)
```

```
x_test_data=StandardScaler( with_mean=False).fit_transform(test_bow)#standardizing the test data
print("shape of x_test after standardization ",x_test_data.shape)
```

```
shape of x_test after bow vectorization (20000, 7677)
shape of x_test after standardization (20000, 7677)
```

```
In [9]: #using time series split method for cross-validation score
from sklearn.model_selection import TimeSeriesSplit
tscv = TimeSeriesSplit(n_splits=10)
from sklearn.linear_model import LogisticRegression
from scipy.stats import uniform
data=[10**(-4),10**(-3),10**(-2),10**(-1),10**0,10**1,10**2,10**3,10**4]#range of hyperparameter
```

```
In [10]: lr=LogisticRegression(penalty='l2',class_weight={1:.5,0:.5})#building logistic regression model
tuned_parameters=[{'C':data}]
```

```
In [10]: #applying the model of logistic regression and using gridsearchcv to find the best hyper parameter
%time
from sklearn.model_selection import GridSearchCV
model = GridSearchCV(lr, tuned_parameters, scoring = 'f1', cv=tscv,n_jobs=-1)#building the gridsearchcv model
model.fit(x_train_data, y_train)#fitting the training data

print(model.best_estimator_)#printing the best_estimator
print(model.score(x_test_data, y_test))#predicting f1 score on test data
```

```
Wall time: 0 ns
LogisticRegression(C=0.001, class_weight={1: 0.5, 0: 0.5}, dual=False,
                  fit_intercept=True, intercept_scaling=1, max_iter=100,
                  multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
                  solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
0.952504987808
```

```
In [11]: lr=LogisticRegression(C=0.001,penalty='l2',class_weight={1:.5,0:.5},n_j
obs=-1)#again building the model to find best hyperparameter
lr.fit(x_train_data,y_train)#fitting the training data
z=lr.decision_function(x_train_data)#checking the signed distance of a
point from hyperplane
print(z)#printing the signed distance
```

```
[ 3.55845076  2.26483878  1.98815951 ...,  7.38021825 -2.68550859
 4.84653495]
```

```
In [12]: wieght_vector=lr.coef_#getting the weight vector
print(wieght_vector.shape)#wieght vector shape
print(wieght_vector[:20])
```

```
(1, 7677)
[[-0.00404381 -0.0055471 -0.00635504 ...,  0.0277215 -0.00865251
  0.00910158]]
```

PERFORMING SPARSITY CHECK WITH L1 REGULARIZATION

```
In [11]: first_20k_points=x_train_data[:20000] #first 20k points
```

```
In [12]: #using time series split method for cross-validation score
from sklearn.model_selection import TimeSeriesSplit
tscv = TimeSeriesSplit(n_splits=10)
from sklearn.linear_model import LogisticRegression
from scipy.stats import uniform
data=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]#range
of hyperparameter
```

```
In [13]: lr=LogisticRegression(penalty='l1',class_weight={1:.5,0:.5})#building l
ogistic regression model
tuned_parameters=[{'C':data}]
```

```
In [14]: #applying the model of logistic regression and using gridsearchcv to find the best hyper parameter
%time
from sklearn.model_selection import GridSearchCV
model = GridSearchCV(lr, tuned_parameters, scoring = 'f1', cv=tscv,n_jobs=-1)#building the gridsearchcv model
model.fit(first_20k_points, y_train[:20000])#fitting the training data

print(model.best_estimator_)#printing the best_estimator
```

```
Wall time: 0 ns
LogisticRegression(C=0.1, class_weight={1: 0.5, 0: 0.5}, dual=False,
                    fit_intercept=True, intercept_scaling=1, max_iter=100,
                    multi_class='ovr', n_jobs=1, penalty='l1', random_state=None,
                    solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
```

```
In [15]: lr=LogisticRegression(C=0.1,penalty='l1',class_weight={1:.5,0:.5},n_jobs=-1)#again building the model to find best hyperparameter
lr.fit(first_20k_points,y_train[:20000])#fitting the training data
```

```
Out[15]: LogisticRegression(C=0.1, class_weight={1: 0.5, 0: 0.5}, dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='ovr', n_jobs=-1, penalty='l1', random_state=None,
                             solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
```

```
In [16]: wieght_vector=lr.coef_#getting the weight vector
print(wieght_vector.shape)#wieght vector shape
print(wieght_vector[:20])
```

```
(1, 7677)
[[ 0.          0.00692072  0.          ...,  0.          0.          0.
    ]]
```

```
In [17]: np.count_nonzero(wieght_vector)
```

```
Out[17]: 1625
```

THUS HERE ONLY 1625 FEATURES ARE NON_ZERO AND REST OF FEATURES WIEGHTS HAVE BECOME ZERO..SPARSITY CHECK IS POSITIVE

PERTUBATION TEST :

AIM: TO CHECK FOR MULTI COLLINEARITY OF FEATURES
STEPS

- 1. GETTING THE WIEGHT VECTOR FROM MODEL AND SAVING IT**
- 2. ADDING NOISE TO THE TRAINING DATA TO GET NEW TRAINING DATA**
- 3. FITTING THE MODEL AGAIN ON NEW DATA**
- 4. GETTING THE WIEGHT VECTOR FROM THIS MODEL 5.ADDING SMALL VALUE TO WEIGHT VECTOR OF BOTH TRAINNG DATA TO REMOVE ANY ERROR**
- 5. FINDING THE PERCENTAGE CHANGE VECTOR**
- 6. GEETING HOW MANY GEATURE HAS CHANGED USING SOME THRESHOLD VALUE(HERE TAKING IT AS 100)**
- 7. PLOTTING THE QUANTILES WITH THIER PERCENTAGE WIGHT VALUE TO CHECK IF COLLINEARITY EXITS OR NOT**

RESULT : TO KNOW WHETHER FEATURES ARE MULTICOLLINEAR OR NOT
AND TO KNOW WHETHER MODEL IS RELIABLE OR NOT #

```
In [19]: #here,we are adding noise to the data  
from scipy.stats import norm  
noise=norm.rvs(size=1)#noise  
x_train_data.data+=noise#adding noise
```

```

In [20]: print('shape of our new train data after adding noise is : ',x_train_d
          ata.shape)#printing shape of new training data

          shape of our new train data after adding noise is : (80000, 7677)

In [21]: #uilding the model using timeSeriesSplit
          from sklearn.model_selection import TimeSeriesSplit
          tscv = TimeSeriesSplit(n_splits=10) # 10 spilts cross validation
          from sklearn.linear_model import LogisticRegression
          from scipy.stats import uniform
          data=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]#value
              range of hyper parameter for grid searchcv
          lr=LogisticRegression(penalty='l2',class_weight={1:.5,0:.5},n_jobs=-1)#
          building the model
          tuned_parameters=[{'C':data}]

In [22]: %time
          from sklearn.model_selection import GridSearchCV
          model = GridSearchCV(lr, tuned_parameters, scoring = 'f1', cv=tscv,n_jo
          bs=-1)#building the gridsearchcv model
          model.fit(x_train_data, y_train)#fiitting the training data

          print('best estimator of our new data is: ',model.best_estimator_)#prin
          ting the best_estimator

          Wall time: 0 ns
          best estimator of our new data is: LogisticRegression(C=0.001, class_w
          eight={1: 0.5, 0: 0.5}, dual=False,
              fit_intercept=True, intercept_scaling=1, max_iter=100,
              multi_class='ovr', n_jobs=-1, penalty='l2', random_state=None,
          e,
              solver='liblinear', tol=0.0001, verbose=0, warm_start=False)

In [23]: # again building the model for finding the wieght vector of the words f
          rom model
          lr=LogisticRegression(C=0.001,penalty='l2',class_weight={1:.5,0:.5},n_j
          obs=-1)#building the logistic regression model
          lr.fit(x_train_data,y_train)#fiting the training model

```

```
new_wieght_vector=lr.coef_  
print(new_wieght_vector.shape)#printing shape of wieght vector  
(1, 7677)
```

```
In [24]: percent_change_vec=np.ones((1,17204))#generating the percent_change_vec  
or to store the percentage change values for each word
```

```
In [25]: wieght_vector=wieght_vector+10**-6 #adding some values to wieght vector  
to avoid error while division  
  
new_wieght_vector=new_wieght_vector+10**-6 #adding some values to wiegh  
t vector to avoid error while division  
  
percent_change_vec=abs((wieght_vector-new_wieght_vector)/wieght_vector)  
*100#calculating the percentage change in the vector
```

```
In [26]: x=(wieght_vector[0][2]-new_wieght_vector[0][2])/wieght_vector[0][2]#jus  
t checking randomly that every value is positive in percent_change_vecto  
r  
print(x)  
  
6124.13518164
```

```
In [27]: print('shape of percent change wieght vector is', percent_change_vec.sh  
ape)#printing shape of percent_change_vector  
  
shape of percent change wieght vector is (1, 7677)
```

```
In [28]: per_change_df=pd.DataFrame(percent_change_vec.T,columns=['CHANGE'])#bui  
lding a dataframe from wieght vector
```

```
In [29]: per_change_df.head()#getting first 5 values
```

Out[29]:

	CHANGE
0	4.114007e+05

	CHANGE
1	1.793059e+02
2	6.124135e+05
3	6.546318e+05
4	1.072202e+06

In [30]: `sorted_Df=per_change_df.sort_values('CHANGE',ascending=True,axis=0)#sorting the dataframe to calculate the quantiles values`
`sorted_Df.describe()#describe function`

Out[30]:

	CHANGE
count	7.677000e+03
mean	8.163010e+05
std	8.943429e+05
min	3.020788e-01
25%	6.567501e+04
50%	5.955684e+05
75%	1.224095e+06
max	1.119868e+07

In [31]: `quantiles=list(i/100 for i in range(0,101,5))#building the list of quantiles value`
`for i in quantiles:`
 `print('sorted_Data {:.2f}th quantiles is {:.3f}'.format(i,sorted_Df['CHANGE'].quantile(i)))#printing the quantiles and thier coreesponding values`

```
sorted_Data 0.00th quantiles is 0.302
sorted_Data 0.05th quantiles is 36.889
sorted_Data 0.10th quantiles is 62.904
sorted_Data 0.15th quantiles is 102.959
sorted_Data 0.20th quantiles is 528.746
sorted_Data 0.25th quantiles is 65675.014
sorted_Data 0.30th quantiles is 180460.231
sorted_Data 0.35th quantiles is 279055.538
sorted_Data 0.40th quantiles is 378386.003
sorted_Data 0.45th quantiles is 482543.855
sorted_Data 0.50th quantiles is 595568.425
sorted_Data 0.55th quantiles is 703414.610
sorted_Data 0.60th quantiles is 811447.629
sorted_Data 0.65th quantiles is 937978.515
sorted_Data 0.70th quantiles is 1076080.247
sorted_Data 0.75th quantiles is 1224094.822
sorted_Data 0.80th quantiles is 1410348.577
sorted_Data 0.85th quantiles is 1629372.501
sorted_Data 0.90th quantiles is 1950475.872
sorted_Data 0.95th quantiles is 2545719.779
sorted_Data 1.00th quantiles is 11198680.280
```

```
In [32]: quantiles=list( i/100 for i in range(95,101,1))#printing the last per
ntiles values because this region is showing abrupt change
percent_change_list=[]#empty percent_change
for i in quantiles:
    print('sorted_Data {:.2f}th quantiles is {:.3f}'.format(i,sorted_Df[
'CHANGE'].quantile(i)))
    percent_change_list.append(sorted_Df['CHANGE'].quantile(i))#building
the list
```

```
sorted_Data 0.95th quantiles is 2545719.779
sorted_Data 0.96th quantiles is 2740753.344
sorted_Data 0.97th quantiles is 2907890.640
sorted_Data 0.98th quantiles is 3243208.220
```


```
sorted_Data 0.99th quantiles is 3921466.304
sorted_Data 1.00th quantiles is 11198680.280
```

```
In [33]: print(percent_change_list)
my_formatted_list = [ '%.2f' % elem for elem in percent_change_list ]#f
ormatted list with string values in it
my_formatted_list=[float(i) for i in my_formatted_list]#formatted list
with float values in it
print(my_formatted_list)#printing formatted list
print(quantiles)#printing quantiles

[2545719.778685397, 2740753.3441018667, 2907890.6404516995, 3243208.219
537627, 3921466.30395979, 11198680.279986566]
[2545719.78, 2740753.34, 2907890.64, 3243208.22, 3921466.3, 11198680.2
8]
[0.95, 0.96, 0.97, 0.98, 0.99, 1.0]
```

```
In [34]: %matplotlib inline
import matplotlib.pyplot as plt
plt.show()
plt.xlabel('QUANTILES VALUES')
plt.ylabel('PERCENTAGE VALUES FOR QUANTILES')
plt.plot(quantiles,my_formatted_list,'r--',label='PERCENTAGE VALUES FO
R QUANTILES')
plt.legend(loc='best')
```

```
Out[34]: <matplotlib.legend.Legend at 0x15e71d05358>
```



FROM THE ABOVE VISUALIZATION , MAIN POINTS ARE:.

1. THAT ONLY 1% OF FEATURES GOT AFFECTED AFTER ADDING NOISE TO THE DATA.
2. VERY LESS COLLINEARITY OF DATA IS PRESENT ,BECAUSE MOST OF THE WEIGHT VECTORS VALUES REMAINS SAME 3.THERFORE, OUR MODEL IS RELIABLE AND WE CAN PROCEED FURTHER TO CHECK ACCURACY ON TEST DATA

CALCULATING THE BEST HYPERPARAMETER ON TRAIN DATA AND CALCULATING THE ACCURACY USING F1-SCORE AND PLOTTING IT

```
In [35]: #using time series split method for cross-validation score
from sklearn.model_selection import TimeSeriesSplit
tscv = TimeSeriesSplit(n_splits=10)
from sklearn.linear_model import LogisticRegression
from scipy.stats import uniform
```

```
data=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]#range  
of hyperparameter
```

```
In [36]: lr=LogisticRegression(penalty='l2',class_weight={1:.5,0:.5},n_jobs=-1)#  
building logistic regression model  
tuned_parameters=[{'C':data}]
```

```
In [37]: #applying the model of logistic regression and using gridsearchcv to fi  
nd the best hyper parameter  
from sklearn.model_selection import GridSearchCV  
model = GridSearchCV(lr, tuned_parameters, scoring = 'f1', cv=tscv,n_jo  
bs=-1)#building the gridsearchcv model  
model.fit(x_train_data, y_train)#fiitting the training data
```

```
Out[37]: GridSearchCV(cv=TimeSeriesSplit(max_train_size=None, n_splits=10),  
error_score='raise',  
estimator=LogisticRegression(C=1.0, class_weight={1: 0.5, 0: 0.  
5}, dual=False,  
fit_intercept=True, intercept_scaling=1, max_iter=100,  
multi_class='ovr', n_jobs=-1, penalty='l2', random_state=Non  
e,  
solver='liblinear', tol=0.0001, verbose=0, warm_start=False),  
fit_params=None, iid=True, n_jobs=-1,  
param_grid=[{'C': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 1  
0000]}],  
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',  
scoring='f1', verbose=0)
```


```
In [38]: results=pd.DataFrame(model.cv_results_)# getting varoius cv_scores and  
train_scores various values of alpha given as parameter and storing it  
in a dataframe  
results#printing the dataframe
```

Out[38]:

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	parar
0	1.371375	0.009375	0.947138	0.957927	0.0001	{'C': 0.000

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	param_alpha
1	2.322609	0.007812	0.956032	0.979998	0.001	{'C': 0.001}
2	4.062930	0.001562	0.954233	0.991260	0.01	{'C': 0.01}
3	7.760476	0.007813	0.946258	0.995863	0.1	{'C': 0.1}
4	14.133621	0.007087	0.937421	0.997608	1	{'C': 1}
5	19.005972	0.009486	0.932540	0.997968	10	{'C': 10}
6	19.856343	0.007924	0.927830	0.998044	100	{'C': 100}
7	26.530732	0.009849	0.925073	0.998082	1000	{'C': 1000}
8	30.367647	0.004799	0.923287	0.998256	10000	{'C': 10000}

9 rows × 31 columns



```
In [39]: %matplotlib inline
import matplotlib.pyplot as plt

mean_test_score=list(results['mean_test_score'])#taking mean_test_score
           values of various alpha into a list
mean_train_score=list(results['mean_train_score'])#taking mean_train_score
           values of various alpha into a list
cv_error_list=[]
train_error_list=[]
for i in mean_test_score:
```

```

i=1-i
i=i*100
cv_error_list.append(i)#appending the list with cv_error
for i in mean_train_score:
    i=1-i
    i=i*100
    train_error_list.append(i)#appending the list with train_error

print(cv_error_list)
C_values_in_10_power=[-4,-3,-2,-1,0,1,2,3,4]#list of alpha values in po
wer of 10
plt.plot(C_values_in_10_power,cv_error_list,label='cv_error')#plotting
alpha with cv_error
plt.plot(C_values_in_10_power,train_error_list,label='train_error')#plo
tting alpha with train_error
plt.xlabel('C value in power of 10 ')
plt.ylabel('cv error and train error')
plt.legend(loc='best')

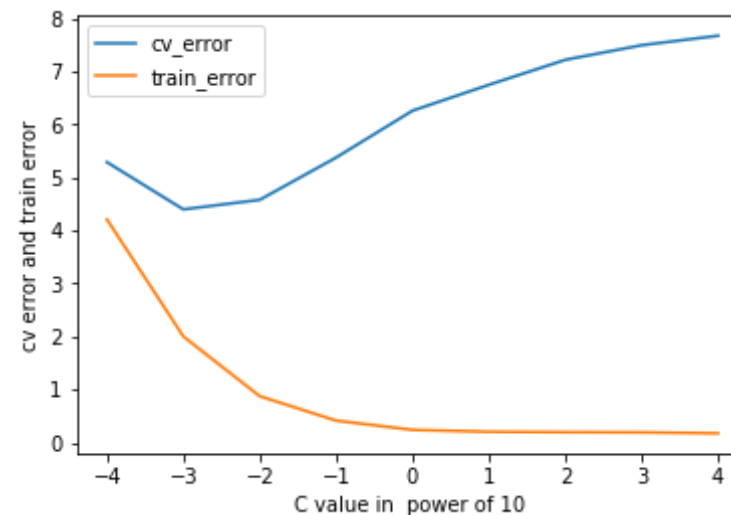
```

```

[5.2862463174302876, 4.3968129672200318, 4.5767022275204798, 5.37416529
13736448, 6.2578765329667103, 6.7460040101474439, 7.2170415778466062,
7.4926855714309131, 7.6713443582323571]

```

Out[39]: <matplotlib.legend.Legend at 0x15e72866b70>



From here, the best hyperparameter value is $c=0.001$ or $\alpha=1000$

NOW GETTING THE TOP 30 FEATURES WORDS FOR POSITIVE AND NEGATIVE WORDS

```
In [40]: #building the model using timeSeriesSplit
from sklearn.model_selection import TimeSeriesSplit
tscv = TimeSeriesSplit(n_splits=10) # 10 splits cross validation
from sklearn.linear_model import LogisticRegression
lr=LogisticRegression(C=0.001,penalty='l2',class_weight={1:.5,0:.5},n_jobs=-1)#building logistic regression model
lr.fit(x_train_data,y_train)
```

```
Out[40]: LogisticRegression(C=0.001, class_weight={1: 0.5, 0: 0.5}, dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='ovr', n_jobs=-1, penalty='l2', random_state=None,
                             solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
```

```
In [41]: z=lr.coef_[0]#getting the weight of the vector
print(z)#printing the weight of the vector

[-0.00411401 -0.00549033 -0.00612414 ...,  0.02776355 -0.00862934
 0.00916463]
```

```
In [42]: a=z.argsort()
print('shape of weight vector is:',a.shape)
top_30_positive=np.take(vectorizer.get_feature_names(),a[17174:])
top_30_negative=np.take(vectorizer.get_feature_names(),a[:30])

shape of weight vector is: (7677,)
```



```
In [43]: print("POSITIVE WORDS\t|\tNEGATIVE WORDS")
for i,j in zip(top_30_positive,top_30_negative):
    print( '{}\t\t|\t\t{}'.format(i,j) )#printing the postive and negat
ive words
```

POSITIVE WORDS | NEGATIVE WORDS

USING BEST HYPERPARAMETER VALUE ON TEST DATA AND PLOTTING THE CONFUSION MATRIX WITH HEATMAP

```
In [44]: #Testing Accuracy on Test data
import seaborn as sns #importing seaborn as sns
from sklearn.metrics import *#importing varoius metrics from sklearn
#building the model
lr.fit(x_test_data,y_test)
y_pred = lr.predict(x_test_data)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100))#printing accuracy
print("Precision on test set: %0.3f"%(precision_score(y_test, y_pred)))
#printing precision score
print("Recall on test set: %0.3f"%(recall_score(y_test, y_pred))) #prin
ting recall
print("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred)))
print("Confusion Matrix of test set:\n [ [TN FP]\n [FN TP] ]\n")
df_cm = pd.DataFrame(confusion_matrix(y_test, y_pred), range(2),range(2)) #generating the heatmap for confusion matrix
sns.set(font scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Accuracy on test set: 96.320%

Precision on test set: 0.961

Recall on test set: 0.998

F1-Score on test set: 0.979


Confusion Matrix of test set:

[[TN FP]

[FN TP]]

IPN IPI I

Out[44]: <matplotlib.axes._subplots.AxesSubplot at 0x15e77d4bef0>



FROM THE ABOVE OBSERVATIONS ,IT IS
FOUND THAT THE BEST HYPERPARAMETER IS
FOUND AS ALPHA=1000 AND IT IS ALSO
HAVING HIGH PRECISION,RECALL VALUE ON
TEST DATA

In [45]: #BOW VECTORIZATION IS COMPLETED FOR LOGISTIC REGRESSION