OBJECTIVE:

- 1. APPLYING LOGISTIC REGRESSION WITH TFIDF WORD2VEC VECTORIZATION
 - PERFORMING PERTUBATION TEST TO CHECK WHETHER OUR DATA FEATURES ARE COLLINER OR NOT AND PLOTTING THE RESULT
 - FINDING THE BEST HYPERPARAMETER USING GRIDSEARCHCV WITH TRAIN DATA AND CROSS-VALIDATION DATA BY PLOTTING THE RESLUTS OF VAROIUS TRAIN DATA AND CROSS VALIDATION DATA
 - USING THE APROPRIATE VALUE OF HYPERPARAMETER, TESTING ACCURACY ON TEST DATA USING F1-SCORE
 - PLOTTING THE CONFUSION MATRIX TO GET THE PRECISOIN ,RECALL VALUE WITH HELP OF HEATMAP

```
In [0]: from sklearn.model_selection import train_test_split
    g the necessary libraries
    from sklearn.model_selection import RandomizedSearchCV
    from sklearn.datasets import *
    from sklearn import naive_bayes
    from sklearn.feature_extraction.text import CountVectorizer
    from sklearn.feature_extraction.text import TfidfVectorizer
    import numpy as np
    import pandas as pd
    from sklearn import *
    from gensim.models import Word2Vec
    import warnings
    warnings.filterwarnings("ignore")
    from tqdm import tqdm
```

In [0]: from google.colab import drive
 drive.mount('/content/gdrive')#geeting the content from the google driv
 e

Drive already mounted at /content/gdrive; to attempt to forcibly remoun

```
t, call drive.mount("/content/gdrive", force remount=True).
In [0]: final processed data=pd.read csv("gdrive/My Drive/final new data.csv")#
        loading the preprocessed data with 100k points into dataframe
In [0]: # getting the counts of 0 and 1 in "SCORE" column to know whether it is
         unbalanced data or not
        count of 1=0
        count of 0=0
        for i in final processed data['Score']:
           if i==1:
            count of 1+=1
           else:
            count of 0+=1
        print(count of 1)
        print(count of 0)
        #it is an imbalanced dataset
        88521
        11479
In [0]: #spliiting the data into train and test data
        x_train,x_test,y_train,y_test=model selection.train test split(final pr
        ocessed data['CleanedText'].values,final processed data['Score'].values
        ,test size=0.2,shuffle=False)
In [0]: # Training my own Word2Vec model using your own text corpus
        list of sent=[]
        for sent in x train:
         list of sent.append(sent.split())#splitting of sentences into words AN
        D appending them to list
        print(x train[0])
        print(list of sent[0])
        word to vector=Word2Vec(list of sent,min count=5,size=100,workers=2)#co
        nstructing my our word to vector
        w t c words=list(word to vector.wv.vocab)
```

```
print("*********
       ******")
       print("sample words ", w t c words[0:20])
       witti littl book make son laugh loud recit car drive along alway sing r
       efrain hes learn whale india droop love new word book introduc silli cl
       assic book will bet son still abl recit memori colleg
       ['witti', 'littl', 'book', 'make', 'son', 'laugh', 'loud', 'recit', 'ca
       r', 'drive', 'along', 'alway', 'sing', 'refrain', 'hes', 'learn', 'whal
       e', 'india', 'droop', 'love', 'new', 'word', 'book', 'introduc', 'sill
       i', 'classic', 'book', 'will', 'bet', 'son', 'still', 'abl', 'recit',
        'memori', 'colleg']
       *******************************
       sample words ['witti', 'littl', 'book', 'make', 'son', 'laugh', 'lou
       d', 'car', 'drive', 'along', 'alway', 'sing', 'refrain', 'hes', 'lear
       n', 'india', 'droop', 'love', 'new', 'word']
#NOW STARTING TF-IDF WEIGHTED WORD-TO-VEC
       model = TfidfVectorizer()
       tf idf matrix = model.fit transform(x train)
       # we are converting a dictionary with word as a key, and the idf as a v
       alue
       dictionary = dict(zip(model.get feature names(), list(model.idf )))
       train tfidf sent vectors =[]# the tfidf-w2v for each sentence/review is
        stored in this list
       for sent in tqdm(list of sent): # for each review/sentence
         sent vec = np.zeros(100) # as word vectors are of zero length
         weight sum =0; # num of words with a valid vector in the sentence/rev
       iew
         for word in sent: # for each word in a review/sentence
          if word in w t c words:
            vec = word to vector.wv[word]
            tf idf = dictionary[word]*(sent.count(word)/len(sent))# dictionary
        [word] = idf value of word in whole courpus
            sent vec += (vec * tf idf)# sent.count(word) = tf valeus of word i
```

```
n this review
           weight sum += tf idf
         if weight sum != 0:
          sent vec /= weight sum
          train tfidf sent vectors.append(sent vec)
       100%| 80000/80000 [06:33<00:00, 203.33it/s]
In [0]: from sklearn.preprocessing import StandardScaler #standarizing the trai
       ning data
       x train data=StandardScaler( with mean=False).fit transform(train tfidf
       sent vectors)
       print(x train data.shape)
       (80000, 100)
In [0]: list of sent=[]
       for sent in x test:
       list of sent.append(sent.split())#splitting of sentences into words AN
       D appending them to list
       print(x test[0])
       print(list of sent[0])
       print('*************
       ***')
       hard find item dont buy mani either came stale got way guick classic no
       netheless
       ['hard', 'find', 'item', 'dont', 'buy', 'mani', 'either', 'came', 'stal
       e', 'got', 'way', 'quick', 'classic', 'nonetheless']
#NOW STARTING TF-IDF WEIGHTED WORD-TO-VEC
       model = TfidfVectorizer()
       model.fit transform(x train)
```

```
model.transform(x test)
        # we are converting a dictionary with word as a key, and the idf as a v
        alue
        dictionary = dict(zip(model.get feature names(), list(model.idf )))
        test tfidf sent vectors =[]# the tfidf-w2v for each sentence/review is
         stored in this list
        for sent in tqdm(list of sent): # for each review/sentence
          sent vec = np.zeros(100) # as word vectors are of zero length
          weight sum =0; # num of words with a valid vector in the sentence/rev
        iew
          for word in sent: # for each word in a review/sentence
           if word in w t c words:
             vec = word to vector.wv[word]
             tf idf = dictionary[word]*(sent.count(word)/len(sent))# dictionary
        [word] = idf value of word in whole courpus
             sent vec += (vec * tf idf)# sent.count(word) = tf valeus of word i
        n this review
             weight sum += tf idf
          if weight sum != 0:
           sent vec /= weight sum
           test tfidf sent vectors.append(sent vec)
                       | 20000/20000 [01:51<00:00, 179.80it/s]
        100%
In [0]: from sklearn.preprocessing import StandardScaler #standarizing the trai
        ning data
        x test data=StandardScaler( with mean=False).fit transform(test tfidf s
        ent vectors)
        print(x test data.shape)
        (20000, 100)
In [0]: #using time series split method for cross-validation score
        from sklearn.model selection import TimeSeriesSplit
        tscv = TimeSeriesSplit(n splits=10)
        from sklearn.linear model import LogisticRegression
        from scipy.stats import uniform
```

```
parameter range=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10
        **4]#range of hyperparameter
In [0]: lr=LogisticRegression(penalty='l2',class weight={1:.5,0:.5},n jobs=-1)#
        building logistic regression model
        tuned parameters=[{'C':parameter range}]
In [0]: #applying the model of logistic regression and using gridsearchev to fi
        nd the best hyper parameter
        %time
        from sklearn.model selection import GridSearchCV
        model = GridSearchCV(lr, tuned parameters, scoring = 'f1', cv=tscv,n jo
        bs=-1)#building the gridsearchcv model
        model.fit(x train data, y train)#fiitting the training data
        print(model.best estimator )#printing the best estimator
        print(model.score(x test data, y test))#predicting f1 score on test da
        CPU times: user 5 \mus, sys: 0 ns, total: 5 \mus
        Wall time: 9.06 µs
        LogisticRegression(C=0.1, class weight={1: 0.5, 0: 0.5}, dual=False,
                  fit intercept=True, intercept scaling=1, max iter=100,
                  multi class='warn', n jobs=-1, penalty='l2', random state=Non
        e,
                  solver='warn', tol=0.0001, verbose=0, warm start=False)
        0.9398098682227883
In [0]: | lr=LogisticRegression(C=0.1, penalty='l2', class weight={1:.5,0:.5}, n job
        s=-1)#building model for getting wieght vector
        lr.fit(x train data,y train)#fitting the training data
        z=lr.decision function(x train data)#checking the signed distance of a
         point from hyperplane
        print(z)#printing the signed distance
        [3.3817815 3.72982275 4.43153667 ... 2.21317907 2.13869536 3.92686093]
In [0]: wieght vector=lr.coef #getting the weight vector
```

```
print(wieght vector.shape)#wieght vector shape
print(wieght vector[:10])
(1, 100)
[[-0.16501958 -0.21457623  0.13722374 -0.11166733  0.22102247 -0.012869
  -0.06977602 0.23810098 0.1583446 -0.01580348 -0.56143422 0.044593
77
  -0.27458457 -0.02144127 0.23681264 0.3060031
                                                0.13795873 -0.069243
  0.05188021 0.02388875 -0.09548089 0.32181222 -0.12674304 0.051333
34
  0.33665299 0.01795062 -0.21469577 0.36660306 -0.02321149 0.405117
09
  -0.35227165 0.26881642 -0.57090952 0.22466264 0.36935362 0.121718
55
  -0.20692414 -0.22261148 -0.23196746 0.11794993 0.0297931 -0.045191
  -0.62683822 0.30744293 0.70980847 -0.26242057 0.25476866 0.133513
17
  -0.11249602 0.59678395 0.07131134 -0.07589831 0.00088051 -0.056469
  -0.11040431 -0.05918636 -0.25525033 -0.20574963 -0.21631142 0.335051
51
  0.10419019 - 0.35587076 - 0.0598645 - 0.20009404 0.19458033 0.117281
63
  0.17198177 - 0.60256641 \ 0.1729033 - 0.10077599 \ 0.17313286 \ 0.090180
91
  -0.31024814 -0.25297867 -0.1085571 -0.02121957 -0.31238821 -0.139704
66
  0.06517518 0.03902697 0.05164876 0.2659251
                                               0.2649679
                                                           0.067296
  -0.24576033 -0.10450755 0.06208229 0.01022223 -0.03161103 -0.051281
43
  -0.04471847 0.00125878 -0.10737445 0.20960195 -0.11948915 0.015933
72
```

PERTUBATION TEST:

AIM: TO CHECK FOR MULTI COLLINEARITY OF FEATURES STEPS

- 1. GETTING THE WIEGHT VECTOR FROM MODEL AND SAVING IT</br>
- 2. ADDING NOISE TO THE TRAINING DATA TO GET NEW TRAINING DATA</br>
- 3. FITTING THE MODEL AGAIN ON NEW DATA</br>
- 4. GETTING THE WIEGHT VECTOR FROM THIS MODEL</br>
 5.ADDING SMALL

 VALUE TO WEIGHT VECTOR OF BOTH TRAINING DATA TO REMOVE ANY ERROR
- 5. FINDING THE PERCENTAGE CHANGE VECTOR
- 6. GEETING HOW MANY GEATURE HAS CHANGED USING SOME THRESHOLD VALUE(HERE TAKING IT AS 100)
- 7. PLOTTING THE QUANTILES WITH THIER PERCENTAGE WIGHT VALUE TO CHECK IF COLLINEARITY EXITS OR NOT</i>

RESULT : TO KNOW WHETHER FEATURES ARE MULTICOLLINEAR OR NOT AND TO KNOW WHETHER MODEL IS RELIABLE OR NOT

```
In [0]: #here,we are adding noise to the data
from scipy.stats import norm
noise=norm.rvs(size=1)#noise
x_train_data.data+=noise#adding noise
```

In [0]: print('shape of our new train data after adding noise is : ',x_train_d
 ata.shape)#printing shape of new training data

shape of our new train data after adding noise is : (80000, 100)

```
In [0]: #uilding the model using timeSeriesSplit
    from sklearn.model_selection import TimeSeriesSplit
    tscv = TimeSeriesSplit(n_splits=10) # 10 spilts cross validation
    from sklearn.linear_model import LogisticRegression
    from scipy.stats import uniform
    data=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]#value
    range of hyper parameter for grid searchcv
    lr=LogisticRegression(penalty='l2',class_weight={1:.5,0:.5},n_jobs=-1)#
```

```
building the model
        tuned parameters=[{'C':data}]
In [0]: %time
        from sklearn.model selection import GridSearchCV
        model = GridSearchCV(lr, tuned parameters, scoring = 'f1', cv=tscv,n jo
        bs=-1)#building the gridsearchcv model
        model.fit(x train data, y train)#fiitting the training data
        print('best estimator of our new data is: ',model.best estimator )#prin
        ting the best estimator
        CPU times: user 4 μs, sys: 0 ns, total: 4 μs
        Wall time: 9.54 µs
        best estimator of our new data is: LogisticRegression(C=0.1, class wei
        ght={1: 0.5, 0: 0.5}, dual=False,
                  fit intercept=True, intercept scaling=1, max iter=100,
                  multi class='warn', n jobs=-1, penalty='l2', random state=Non
        e,
                  solver='warn', tol=0.0001, verbose=0, warm start=False)
In [0]: # again building the model for finding the wieght vector of the words f
        rom model
        lr=LogisticRegression(C=10, penalty='l2', class weight={1:.5,0:.5}, n jobs
        =-1)#building the logistic regression model
        lr.fit(x train data,y train)#fiting the training model
        new wieght vector=lr.coef
        print(new wieght vector.shape)#printing shape of wieght vector
        (1, 100)
In [0]: percent change vec=np.ones((1,100))#generating the percent change vetor
         to store the percentage change values for each word
In [0]: wieght vector=wieght vector+10**-6 #adding some values to wieght vector
         to avoid error while division
        new wieght vector=new wieght vector+10**-6 #adding some values to wiegh
```

```
t vector to avoid error while division
        percent change vec=abs((wieght vector-new wieght vector)/wieght vector)
        *100#calculating the percentage change in the vector
In [0]: x=abs((wieght vector[0][2]-new wieght vector[0][2])/wieght vector[0][2
        ])#just checking randomly that every value is positive in percent change
        vector
        print(x)
        0.00283117933505747
In [0]: print('shape of percent change wieght vector is', percent change vec.sh
        ape)#printing shape of percent change vector
        shape of percent change wieght vector is (1, 100)
In [0]: per change df=pd.DataFrame(percent change vec.T,columns=['CHANGE'])#bui
        lding a dataframe from wight vector
        per change df.head()#getting first 5 values
Out[0]:
           CHANGE
        0 8.252026
         1 5.562601
        2 0.283118
         3 6.209997
         4 1.448420
In [0]: sorted Df=per change df.sort values('CHANGE', ascending=True, axis=0)#sor
        ting the dataframe to calculate the quantiles values
        sorted Df.describe()#describe function
```

```
Out[0]: CHANGE

count 100.000000

mean 21.766990

std 96.913322

min 0.041925

25% 2.558892

50% 5.629969

75% 8.782768
```

max

947.544153

```
In [0]: quantiles=list( i/100 for i in range(0,101,5))#building the list of qua
        ntiles value
        for i in quantiles:
          print('sorted Data {:.2f}th quantiles is {:7.3f}'.format(i,sorted Df[
        'CHANGE'].quantile(i)))#printing the quantiles and thier coreesponding
         values
        sorted Data 0.00th quantiles is
                                           0.042
        sorted Data 0.05th quantiles is
                                          0.446
        sorted Data 0.10th quantiles is
                                           1.014
        sorted Data 0.15th quantiles is
                                          1.443
        sorted Data 0.20th quantiles is
                                          1.793
        sorted Data 0.25th quantiles is
                                          2.559
        sorted Data 0.30th quantiles is
                                          3.411
        sorted Data 0.35th quantiles is
                                          3.770
        sorted Data 0.40th quantiles is
                                          4.985
        sorted Data 0.45th quantiles is
                                          5.261
        sorted Data 0.50th quantiles is
                                          5.630
```

```
sorted Data 0.55th quantiles is
                                          6.221
                                          6.526
        sorted Data 0.60th quantiles is
        sorted Data 0.65th quantiles is
                                          7.206
        sorted Data 0.70th quantiles is
                                          8.266
        sorted Data 0.75th quantiles is
                                          8.783
        sorted Data 0.80th quantiles is
                                          9.965
        sorted Data 0.85th quantiles is 14.053
        sorted Data 0.90th quantiles is 27.903
        sorted Data 0.95th quantiles is 57.288
        sorted Data 1.00th quantiles is 947.544
In [0]: quantiles=list( i/100 for i in range(95,101,1))#printing the last perce
        ntiles values because this region is showing abrupt change
        percent change list=[]#empty percent change
        for i in quantiles:
          print('sorted Data {:.2f}th quantiles is {:7.3f}'.format(i,sorted Df[
        'CHANGE'].quantile(i)))
          percent change list.append(sorted Df['CHANGE'].quantile(i))#building
         the list
        sorted Data 0.95th quantiles is 57.288
        sorted Data 0.96th quantiles is 64.478
        sorted Data 0.97th quantiles is 79.859
        sorted Data 0.98th quantiles is 159.585
        sorted Data 0.99th quantiles is 174.877
        sorted Data 1.00th quantiles is 947.544
In [0]: print(percent change list)
        my formatted list = [ '\%.2f' \% ] elem for elem in percent change list ]#f
        ormatted list with string values in it
        my formatted list=[float(i) for i in my formatted list]#formatted list
         with flaot values in it
        print(my formatted list)#printing formatted list
        print(quantiles)#printing quantiles
        [57.28808961407178, 64.47832543785151, 79.85944211906845, 159.584934819
        49412, 174.8767261023266, 947.5441530541704]
        [57.29, 64.48, 79.86, 159.58, 174.88, 947.54]
        [0.95, 0.96, 0.97, 0.98, 0.99, 1.0]
```

```
In [0]: %matplotlib inline
   import matplotlib.pyplot as plt
   plt.show()
   plt.xlabel('QUANTILES VALUES')
   plt.ylabel('PERCENTAGE VALUES FOR QUANTILES')
   plt.plot(quantiles,my_formatted_list,'r--' ,label='PERCENTAGE VALUES FO
   R QUANTILES')
   plt.legend(loc='best')
```

Out[0]: <matplotlib.legend.Legend at 0x7fa1a6a6d978>



FROM THE ABOVE VISUALIZATION, MAIN POINTS ARE:.

- 1. THAT ONLY 1% OF FEATURES GOT AFFECTED AFTER ADDING NOISE TO THE DATA.
- 2. VERY LESS COLLINEARITY OF DATA IS PRESENT ,BECAUSE MOST OF THE WEIGHT VECTORS VALUES REMAINS SAME
 3.THERFORE, OUR MODEL IS RELIABLE AND WE CAN PROCEED FURTHER TO CHECK ACCURACY ON TEST DATA

CALCULATING THE BEST HYPERPARAMETER ON TRAIN DATA AND CALCULATING THE ACCURACY USING F1-SCORE AND PLOTTING IT

```
In [0]: #using time series split method for cross-validation score
        from sklearn.model selection import TimeSeriesSplit
        tscv = TimeSeriesSplit(n splits=10)
        from sklearn.linear model import LogisticRegression
        from scipy.stats import uniform
        data=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]#range
        of hyperparameter
In [0]: | lr=LogisticRegression(penalty='l2',class weight={1:.5,0:.5},n jobs=-1)#
        building logistic regression model
        tuned parameters=[{'C':data}]
In [0]: #applying the model of logistic regression and using gridsearchev to fi
        nd the best hyper parameter
        from sklearn.model selection import GridSearchCV
        model = GridSearchCV(lr, tuned parameters, scoring = 'f1', cv=tscv,n jo
        bs=-1)#building the gridsearchcv model
        model.fit(x train data, y train)#fiitting the training data
Out[0]: GridSearchCV(cv=TimeSeriesSplit(max train size=None, n splits=10),
              error score='raise-deprecating',
              estimator=LogisticRegression(C=1.0, class weight={1: 0.5, 0: 0.
        5}, dual=False,
                 fit intercept=True, intercept scaling=1, max iter=100,
                 multi class='warn', n jobs=-1, penalty='l2', random state=Non
        e,
                 solver='warn', tol=0.0001, verbose=0, warm start=False),
              fit params=None, iid='warn', n jobs=-1,
              00001}1,
              pre dispatch='2*n jobs', refit=True, return train score='warn',
              scoring='f1', verbose=0)
```

Out[0]:

		_		_	_	
	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	parar
0	1.996376	0.014547	0.941283	0.941648	0.0001	{'C': 0.000
1	3.171960	0.014771	0.943710	0.945046	0.001	{'C': 0.001
2	4.698934	0.015680	0.946472	0.950018	0.01	{'C': 0.01}
3	7.046733	0.015911	0.946884	0.950443	0.1	{'C': 0.1}
4	7.538299	0.014789	0.946434	0.950578	1	{'C': '
5	7.361695	0.012899	0.946418	0.950613	10	{'C': 10}
6	7.598320	0.013064	0.946426	0.950609	100	{'C': 100}
7	7.546490	0.016040	0.946426	0.950609	1000	{'C': 1000]

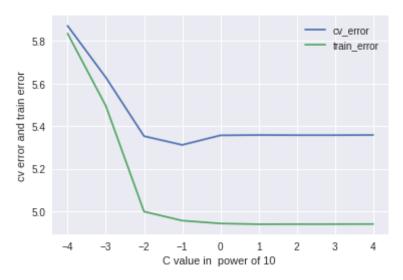
	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	parar
8	7.403473	0.012514	0.946418	0.950605	10000	{'C':

9 rows × 31 columns

In [0]: %matplotlib inline import matplotlib.pyplot as plt mean test score=list(results['mean test score'])#taking mean test score values of various alpha into a list mean train score=list(results['mean train score'])#taking mean train sc ore values of varoius alpha into a list cv error list=[] train error list=[] for i in mean test score: i=1-ii=i*100 cv error list.append(i)#appending the list with cv error for i in mean train score: i=1-i i=i*100train error list.append(i)#appending the list with train error print(cv error list) C values in 10 power=[-4,-3,-2,-1,0,1,2,3,4]#list of alpha values in po wer of 10 plt.plot(C values in 10 power,cv error list,label='cv error')#plotting alpha with cv error plt.plot(C values in 10 power,train error list,label='train error')#plo tting aplhawith train error plt.xlabel('C value in power of 10 ') plt.ylabel('cv error and train error') plt.legend(loc='best') [5.8716940632852825, 5.629002786920556, 5.352804983977344, 5.3116281652

48572, 5.35657351955895, 5.358155325677227, 5.3574406314689575, 5.35744 06314689575, 5.358235102558151]

Out[0]: <matplotlib.legend.Legend at 0x7fa1a630ee48>

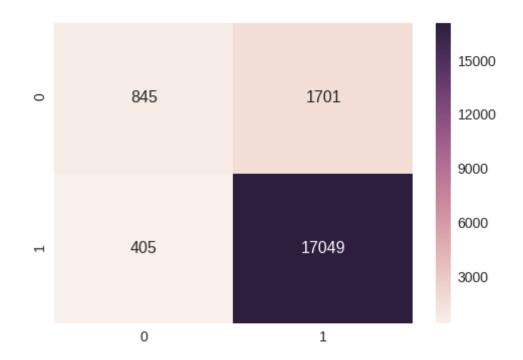


From here, the best hyperparameter value is c=10 or alpha=0.1

USING BEST HYPERPARAMETER VALUE ON TEST DATA AND PLOTTING THE CONFUSION MATRIX WITH HEATMAP

```
In [0]: #Testing Accuracy on Test data
import seaborn as sns #importing seaborn as sns
from sklearn.metrics import *#importing varoius metrics from sklearn
lr=LogisticRegression(C=10,penalty='l2',class_weight={1:.5,0:.5},n_jobs
=-1)#building logistic regression model#building the model
lr.fit(x_test_data,y_test)
y_pred = lr.predict(x_test_data)
```

```
print("Accuracy on test set: %0.3f%"%(accuracy score(y test, y pred)*1
        00))#printing accuracy
        print("Precision on test set: %0.3f"%(precision score(y test, y pred)))
        #printing precision score
        print("Recall on test set: %0.3f"%(recall score(y_test, y_pred))) #prin
        ting recall
        print("F1-Score on test set: %0.3f"%(f1 score(y test, y pred)))
        print("Confusion Matrix of test set:\n [ [TN FP]\n [FN TP] ]\n")
        df cm = pd.DataFrame(confusion matrix(y test, y pred), range(2),range(2)
        )) #generating the heatmap for confusion matrix
        sns.set(font scale=1.4)#for label size
        sns.heatmap(df cm, annot=True,annot kws={"size": 16}, fmt='g')
        Accuracy on test set: 89.470%
        Precision on test set: 0.909
        Recall on test set: 0.977
        F1-Score on test set: 0.942
        Confusion Matrix of test set:
         [ [TN FP]
         [FN TP] ]
Out[0]: <matplotlib.axes. subplots.AxesSubplot at 0x7fa1a69a4208>
```



FROM THE ABOVE OBSERVATIONS, IT IS FOUND THAT THE BEST HYPERPARAMETER IS FOUND AS APLHA=0.1 AND IT IS ALSO HAVING HIGH PRECISION, RECALL VALUE ON TEST DATA

In [0]: #TFIDF WIEGHTED WORD2VEC VECTORIZATION IS COMPLETED FOR LOGISTIC REGRES SION

SUMMARIZING THE VAROIUS VECTORIZATIONS FOR F1-

SCORE, RECALL, PRECISION, ACCURACY ON TEST DATA WITH HELP OF TABLE

```
In [3]: from tabulate import tabulate
In [4]: table = [["BOW",1000,90,99,95],["TF-IDF",1000,90,99,94], ["AVG WORD 2 V
        EC", 0.1, 91, 97, 94], ["TFIDF AVG WORD 2 VEC", 0.1, 91, 98, 94]]
       print (tabulate(table))
       BOW
                            1000
                                    90 99 95
       TF-IDF
                                    90 99 94
                            1000
       AVG WORD 2 VEC
                               0.1 91 97 94
       TFIDF AVG WORD 2 VEC
                               0.1 91 98 94
       headers=['VECTORIZATION','ALPHA','PRECISION','RECALL','F1 SCORE']
In [5]:
        print (tabulate(table, headers, tablefmt="fancy grid"))
                                 ALPHA |
                                           PRECISION |
                                                        RECALL
                                                                  F1 SCORE
         VECTORIZATION
                                1000
                                                 90
                                                            99
         BOW
                                                                        95
         TF-IDF
                                1000
                                                 90
                                                            99
                                                                        94
                                   0.1
                                                 91
                                                            97
         AVG WORD 2 VEC
                                                                        94
```

